

**A Rubric-Based Framework for Search Advertising Quality Evaluation: Theoretical
Foundations and Applied Case Studies**

Diego Palencia Robles

MSc, Traffic & Digital Marketing (IEBS – UCAM)

PhD Candidate, High-Performance Management (Universidad Galileo)

Prepared in application for: AI Google Ads Digital Practitioner, TELUS

Guatemala City, Guatemala · July 2026

Abstract

This paper presents a rubric-based framework for evaluating the relevance and quality of search advertisements, developed in direct response to the practical demands of AI-assisted ad-quality evaluation programs. Grounded in established search-intent theory (Broder, 2002; Jansen et al., 2008), Google's published Quality Score methodology, and inter-rater reliability principles (Cohen, 1960), the framework operationalizes ad evaluation as a reproducible measurement task rather than a subjective judgment. Three applied case studies demonstrate the framework in practice: (a) a five-dimension relevance evaluation of nine live advertisements across three search queries, (b) a complete Google Ads campaign architecture built around a single conversion objective, and (c) a thirty-keyword intent-classification study for a services market. Across all three studies, evaluation outcomes are expressed as documented, re-derivable scores rather than holistic impressions, consistent with the calibration standard required of multi-rater evaluation programs. The paper concludes by mapping each documented competency against the requirements of the AI Google Ads Digital Practitioner role, arguing that structured, theory-informed evaluation is directly transferable to the human-in-the-loop assessment of AI-powered advertising systems.

Keywords: ad quality evaluation, search intent, Quality Score, keyword match types, inter-rater reliability, decision intelligence, Google Ads

Introduction

Search advertising sits at the intersection of two evaluation problems. The first is technical: does an advertisement match the literal keywords and structural signals of a query? The second is interpretive: does the advertisement satisfy the need that motivated the search in the first place? Google's own Quality Score methodology treats this as a measurable diagnostic rather than a matter of opinion, built from three components (expected click-through rate, ad relevance, and landing page experience) each rated relative to comparable advertisers (Google, n.d.-b). The same logic increasingly extends beyond automated scoring: as AI systems take on more of the work of ranking and serving advertisements, human evaluators are needed to judge, calibrate, and correct those systems' outputs against real search intent.

This is precisely the applied problem addressed by AI-assisted advertising evaluation roles such as the AI Google Ads Digital Practitioner position at TELUS, which requires assessing ad relevance against user search behavior, keywords, and intent, and producing feedback that can improve AI-powered advertising systems. That responsibility is not well served by holistic, unstated judgment. It requires a rubric: a small number of named dimensions, each scored independently, each defensible to a second reviewer working from the same evidence.

This paper develops such a rubric and demonstrates it across three applied case studies. Section 2 reviews the theoretical concepts that inform the framework, search-intent classification, Quality Score methodology, keyword match-type architecture, inter-rater reliability, and decision intelligence as an applied discipline. Section 3 describes the evaluation methodology used across the three studies. Section 4 presents the case studies themselves: an ad-quality evaluation, a full campaign blueprint, and a keyword intent-mapping study, each

conducted independently but unified by the same evidentiary standard. Section 5 discusses how the resulting body of evidence maps onto the specific requirements of AI-assisted ad-evaluation work, and Section 6 concludes.

Theoretical Framework and Key Concepts

Search Intent Classification

Broder (2002) proposed the foundational taxonomy still used across the search industry, classifying web queries as informational (the user wants to learn something), navigational (the user wants to reach a specific site), or transactional (the user wants to complete an action, such as a purchase or download). Jansen et al. (2008) extended this taxonomy into an automatically classifiable system, demonstrating that intent could be inferred at scale from query structure and log behavior rather than requiring manual annotation of every search. Search-engine-marketing practice has since added a fourth, practitioner-defined category (commercial investigation) describing queries where a user is comparing options ahead of a purchase decision without yet being ready to transact. All four categories appear directly in the case studies below: Case Study I evaluates an advertisement against a query explicitly labeled commercial investigation, and Case Study III applies a four-part classification (informational, commercial, transactional, navigational) to a full keyword universe.

Ad Relevance and Quality Score

Google's Quality Score is a diagnostic, not a ranking input, calculated on a 1–10 scale from three components: expected click-through rate, ad relevance, and landing page experience, each benchmarked against other advertisers competing for the same query over a rolling window (Google, n.d.-b). The diagnostic function of this metric (telling an advertiser where an ad or

landing page is underperforming relative to intent, rather than simply scoring it) is the direct model for the five-dimension rubric used in Case Study I and the per-component diagnostic applied in Case Study II.

Keyword Match-Type Architecture

Keyword match types govern how closely a search query must correspond to an advertiser's keyword before an advertisement becomes eligible for the auction. Exact match offers the narrowest, most controlled reach; phrase match requires the keyword's meaning to be present within a broader query; broad match (the default type) extends reach to related searches beyond the literal keyword (Google, n.d.-a). Because broader match types subsume narrower ones, match-type selection is fundamentally a precision–reach trade-off, and it is one of the structural decisions evaluated directly in Case Studies II and III.

Inter-Rater Reliability and Evaluation Calibration

Any evaluation process that depends on human judgment (whether a single rater scoring advertisements or a distributed team calibrating against shared guidelines) inherits a basic measurement problem: judgments must agree beyond what chance alone would produce. Cohen (1960) formalized this with a coefficient of agreement for nominal-scale judgments, still the standard reference point for inter-rater reliability across fields that rely on trained human evaluators. The relevance of this concept to AI-assisted advertising evaluation is direct: a program built around calibration sessions and consensus discussions is, in effect, running a reliability check on its raters, and a rubric with named, independently scored dimensions (rather than a single holistic verdict) is what makes that check possible in the first place.

Applied Decision Intelligence as a Framework

Decision intelligence, as defined by Kozyrkov (2019), is the discipline of turning information into better action at scale, combining applied data science, social science, and managerial science into a single practice organized around the decision itself rather than the dataset. Applied to advertising evaluation, this reframes each ad-relevance verdict, campaign-structure choice, and keyword classification not as an isolated task but as a decision, made under uncertainty, that should be documented well enough to be reviewed, challenged, and improved. This is the organizing principle behind the case studies that follow internally referred to by the reviewer as DJPR (Decision Journey Process Reduction), an applied naming convention for the same underlying discipline.

Methodology

Three independent studies were conducted, unified by a single evidentiary standard: every judgment is expressed as a named, scored dimension rather than a holistic rating, and every score is accompanied by a written justification sufficient for a second reviewer to reproduce or challenge it.

Case Study I applies a five-dimension rubric: (1) keyword alignment, (2) intent match, (3) ad relevance, (4) landing page fit, and (5) trust signals — to real, live advertisements returned for real search queries. Each dimension is scored on the same 1–10 scale used by Google’s own Quality Score diagnostic (Google, n.d.-b), chosen for direct comparability with industry-standard evaluation rather than an arbitrary scale.

Case Study II applies campaign-architecture methodology: ad groups organized by search intent (Broder, 2002), keyword match-type selection informed by the precision–reach trade-off described in Section 2.3, and a Quality-Score-style diagnostic applied to the resulting structure.

Case Study III applies intent-classification methodology, extending Broder’s (2002) three-part taxonomy with the practitioner-defined commercial-investigation category, and prioritizing a keyword universe by proximity to conversion rather than by raw search volume alone.

All three studies were conducted by a single evaluator; the limitations of that design, and its relationship to formal inter-rater reliability, are addressed in Section 5.

Applied Case Studies

Case Study I: Ad Quality Evaluation

This study evaluated nine live advertisements across three real search queries using the five-dimension rubric described in Section 3. Table 1 defines each rubric dimension.

Table 1

Five-Dimension Ad Relevance Rubric


Dimension	What It Measures
Keyword Alignment	Correspondence between the query’s literal terms and the advertisement’s copy
Intent Match	Whether the advertisement addresses the underlying reason for the search, not just its literal wording
Ad Relevance	Clarity, specificity, and quality of the advertisement copy itself
Landing Page Fit	Correspondence between the advertisement’s promise and the page it links to
Trust Signals	Credibility markers (social proof, domain authority, specificity) supporting the advertisement’s claims

Note. Each dimension is scored on a 1–10 scale, consistent with Google’s published Quality Score methodology (Google, n.d.-b).

Query 1: *“affordable CRM software for small business”* (commercial investigation intent). Three advertisements were evaluated. An advertisement for HubSpot CRM scored 9.4 of 10 overall, addressing both query modifiers (affordable and small business) directly in its headline through a free-tier offer, reinforced by a landing-page path and a social-proof claim specific to the small-business segment. An advertisement for Salesforce scored 5.8: strong on trust signals through established brand recognition, but explicitly signaling an enterprise audience (“Enterprise-level features,” “Top Companies”), which contradicts the price-sensitive, small-business intent behind the query. A third advertisement, from a general software-comparison site, scored 3.0, never mentioning the product category (“CRM”) and offering minimal trust signals.

Query 2: *“online Spanish courses for adults”* (transactional intent). An advertisement for Babbel scored 9.3, directly addressing both the audience (adults) and modality (online) while reframing the adult learner’s core constraint (limited time) as a feature (“15 Min/Day”). A second advertisement, for a conventional language school, scored 6.1: competently written, but omitting any reference to the adult audience or a confirmed online modality leading to a specific, documented recommendation that the advertiser add both terms to its headline. Figure 1 presents the overall scores across both queries.

Across both queries, the pattern was consistent: the lowest-scoring advertisements were not poorly written in isolation they were well-written for a different audience or a different intent than the one the query expressed. This distinction between copy quality and intent match is precisely what a single holistic relevance judgment would fail to separate, and precisely what a five-dimension rubric is designed to expose.

Figure 1*Overall Relevance Scores by Evaluated Advertisement*

Note. Scores reflect the mean of five rubric dimensions per advertisement (Table 1), rounded to one decimal place.

Case Study II: Search Campaign Blueprint

This study designed a complete Google Ads account architecture for LexGuatemala, an immigration and corporate law firm, organized around a single key performance indicator: cost per consultation lead. The architecture comprises four ad groups (Table 2), each aligned to a distinct search intent following the taxonomy in Section 2.1, together totaling 24 keywords and 8 responsive search ad assets.


Table 2*Ad Group Architecture, LexGuatemala Campaign*

Ad Group	Search Intent	Representative Keyword	Match Type
AG-01 Visa & Immigration Consultation	Transactional	[abogado de visas Guatemala]	Exact
AG-02 Corporate & Business Law	Commercial Investigation	[abogado corporativo Guatemala]	Exact
AG-03 US Diaspora (English)	Transactional	[Guatemala immigration lawyer]	Exact
AG-04 Brand & Competitor Defense	Navigational	[LexGuatemala]	Exact

Note. All four ad groups maintain independent negative-keyword lists to prevent cross-group query overlap.

Ad group AG-03 replicates the transactional strategy of AG-01 for an English-speaking diaspora audience, with independently written copy rather than direct translation for example, “Guatemala Immigration Lawyer · Online Consultations · Trusted Legal Experts · Get Help Now” and its own negative-keyword list, built specifically to exclude U.S.-domestic immigration queries that would otherwise be structurally ambiguous with the target query set. This bilingual design is direct, demonstrated evidence of working proficiency in both English and Spanish within a single applied campaign.

A Quality-Score-style diagnostic (Section 2.2) was applied to each ad group. For AG-01, both expected click-through rate and ad relevance rated above average relative to comparable advertisers, while landing page experience rated only average the diagnosed constraint on the group’s overall score (Figure 2). The corresponding, documented recommendation was landing-page-specific: incorporate the primary keyword into the page’s main heading, reduce page load time, and surface a trust signal above the fold, rather than adjusting advertisement copy that was already performing above benchmark.

Figure 2*Quality Score Component Diagnostic, Ad Group 01*

Note. Diagnostic ratings are relative to other advertisers competing for the same queries over a rolling 90-day window, consistent with Google's Quality Score methodology (Google, n.d.-b).

Case Study III: Keyword Intent Mapping

This study classified a thirty-keyword universe for an online fitness-coaching service in the Guatemalan market using the four-part intent taxonomy described in Section 2.1 (Table 3), and prioritized the resulting keyword set by proximity to conversion rather than by search volume alone.


Table 3*Intent Classification Framework Applied to the Keyword Universe*

Intent Category	Defining Characteristic	Budget Prioritization
Informational	Seeking knowledge; not yet evaluating providers	Low (nurture rather than convert)
Commercial Investigation	Comparing providers ahead of a decision	Medium (converts within 2–3 days of first contact)
Transactional	Ready to act	High (prioritized for budget allocation)
Navigational	Searching for a known brand by name	Protect, not conquest

Five keywords were classified as top-priority transactional queries (for example, “*entrenador personal online Guatemala*” and “*rutina personalizada online precio*,” exact- and phrase-matched respectively) while informational queries such as “*cómo empezar a hacer ejercicio en casa*” were retained for reach but assigned the lowest bidding priority. Figure 3 presents the resulting share of search volume across the four intent stages, calculated from the classified keyword set: volume falls from 100% at the informational stage to 18% at the navigational stage, with the steepest single decline occurring between the commercial-investigation and transactional stages. The budget implication is direct the highest-value inflection point in the funnel is not where search volume is largest, but where it drops fastest.

Figure 3

Share of Search Volume by Intent Stage



Note. Values represent the proportion of total classified search volume retained at each successive intent stage.

Discussion

The three case studies were conducted independently, for three different markets and objectives, but they share a single methodological commitment: every evaluative judgment (an advertisement's relevance, a campaign's structure, a keyword's priority) is expressed as a named, scored, and justified decision rather than an unstated impression. Table 4 maps that body of evidence directly against the stated requirements of the AI Google Ads Digital Practitioner role.

Table 4

Alignment Between Position Requirements and Demonstrated Evidence

Requirement	Evidence	Source
Evaluate ad relevance against search queries, keywords, and intent	Nine advertisements scored across three queries on one five-dimension rubric; two advertisements on the same query scored 9.4 and 3.0, with the gap fully documented	Case Study I
Search intent	Four-part intent framework applied to a 30-keyword universe with priority ranking	Case Study III
Keyword targeting and match types	24 keywords architected across Exact, Phrase, and Broad match in a live campaign; 30 more classified by intent and match type	Case Studies II–III
Quality Score	Per-ad-group diagnostic (expected CTR, ad relevance, landing page experience) with a specific optimization action attached	Case Study II
Landing page experience	Scored as one of five rubric dimensions on every evaluated advertisement; rated a second time within the Quality Score diagnostic	Case Studies I–II
Campaign optimization; conversion-focused advertising	Full four-ad-group architecture built around one stated KPI cost per consultation lead	Case Study II
Apply guidelines consistently; attention to detail	Identical rubric applied without exception across all nine evaluated advertisements	Case Study I
Provide clear, actionable feedback	Every evaluation closes with a specific, written recommendation, down to the exact headline revision proposed	Case Study I

One limitation of the current body of work is that all three evaluations were conducted by a single rater. Cohen's (1960) coefficient of agreement (and the broader discipline of inter-rater reliability it founded) exists precisely to quantify whether a second, independent evaluator would reach the same verdict from the same evidence, which single-rater work cannot itself demonstrate. This is not a limitation the case studies overlook: the five-dimension rubric in Case Study I and the intent taxonomy in Case Study III were both built with named, discrete dimensions specifically so that a second rater could apply them and a coefficient of agreement could, in principle, be computed. Calibration in a multi-rater setting (the explicit responsibility described in the target role) is the natural next application of a framework built this way, rather than a departure from it.

Conclusion

This paper set out to do two things: ground search-advertising evaluation in established theory, and demonstrate that grounding through applied work rather than assertion. Search-intent classification (Broder, 2002; Jansen et al., 2008), Quality Score methodology (Google, n.d.-b), keyword match-type architecture (Google, n.d.-a), inter-rater reliability (Cohen, 1960), and decision intelligence as an applied discipline (Kozyrkov, 2019) are not decorative citations here each maps onto a specific, documented decision inside the three case studies. The resulting evidence base (nine scored advertisements, a complete four-ad-group campaign architecture, and a thirty-keyword intent classification) is offered as a direct answer to what AI-assisted advertising evaluation work actually requires: not a general marketing background, but the specific discipline of turning a search query into a defensible, reproducible judgment about the advertisement meant to answer it.

References

- Broder, A. (2002). A taxonomy of web search. *ACM SIGIR Forum*, 36(2), 3–10.
<https://doi.org/10.1145/792550.792552>
- Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46. <https://doi.org/10.1177/001316446002000104>
- Google. (n.d.-a). About keyword matching options. *Google Ads Help*. Retrieved July 27, 2026, from <https://support.google.com/google-ads/answer/7478529>
- Google. (n.d.-b). About Quality Score for Search campaigns. *Google Ads Help*. Retrieved July 27, 2026, from <https://support.google.com/google-ads/answer/6167118>
- Jansen, B. J., Booth, D. L., & Spink, A. (2008). Determining the informational, navigational, and transactional intent of Web queries. *Information Processing & Management*, 44(3), 1251–1266. <https://doi.org/10.1016/j.ipm.2007.07.015>
- Kozyrkov, C. (2019, August 2). Introduction to decision intelligence: A new discipline for leadership in the AI era. *Medium*.
<https://medium.com/data-science/introduction-to-decision-intelligence-5d147ddab767>